

ServiceNow Impact

Application overview, business case and activation — Australia release

Prepared from official ServiceNow Australia-release Impact documentation

Audience: executive sponsors, IT leadership, platform owners and architects

Agenda

Eight questions this briefing answers

- 1. What ServiceNow Impact actually is — and what it is not
- 2. Why ServiceNow built it, and which problems it solves
- 3. The architecture: what sits where, and what data moves
- 4. The feature catalogue by pillar
- 5. What it looks like in practice — worked client scenarios
- 6. Business benefits, by stakeholder and by category
- 7. How it differs from traditional support
- 8. Prerequisites, activation steps and the first 90 days

Executive summary

Three things to remember about ServiceNow Impact

1 · Prove value

Objectives → outcomes → success metrics → benchmarks → monetised Value Reports with visible assumptions.

2 · Know your health

HealthScan across manageability, performance, security, upgradeability and user experience; Scan Engine findings with real-time prevention and AI fixes; off-instance monitoring in Instance Observer.

3 · Plan and execute

Capability maps show what you own and use; Product Adoption Roadmaps sequence what comes next; Accelerators supply fixed-scope expert capacity.

1. What ServiceNow Impact is

A premium customer success product — not a normal application

- Impact is a subscription-based customer success product that combines AI-powered tools with hands-on ServiceNow expertise.
- It is delivered in three parts:
 - A digital workspace inside our own instance — the Impact Store Application
 - A ServiceNow-hosted expert space — the Impact Delivery Instance (IDI)
 - Human expertise and service entitlements — the Impact Squad, Developer Support, learning credits and Accelerators
- Stated core functions: plan implementation, track product adoption, maintain platform health, resolve platform issues, access expert support.
- From Yokohama onward the Impact Store Application is the exclusive hub for new Impact features — the IDI does not receive new features.
- It does not replace platform administration, change control or Now Support case handling.

Why it exists and what it solves

From ticket resolution to value realisation

The shift

Ticket-based support answers “something is broken”. Impact answers “are we getting value, is it healthy, and what next?”

Client problem	What Impact provides	Where it lives
Cannot prove the value of the investment	Objectives → measurable outcomes → benchmark/goal trend → monetised Value Reports	Value Management
Bought modules nobody uses	Capability Maps, Product Adoption Roadmaps, entitlement visibility	Product Adoption
Hidden technical debt	HealthScan definitions, Scan Engine findings, Real-Time Prevention, AI fixes	Platform Health
Problems only found at upgrade time	Upgrade readiness, update-set scanning, Proactive Code Check	Platform Health
Performance found by users, not us	Off-instance monitoring, alerts, triage, root-cause correlation	Instance Observer
'What should we do next?'	Recommendation engine, squad guidance, roadmaps, Accelerators	Impact home / IDI
Expert capacity is the bottleneck	Developer Support, Accelerators, Initiatives, learning credits	Experts & learning

Where Impact fits on the platform

Around and above the AI Platform — it does not replace our processes

Your instance

Impact Store Application ·
Platform Health · Value
Management packs ·
Instance Observer telemetry

Service Bridge

Secure, bi-directional sync —
status must read Active
replication

Impact Delivery Instance

Impact Squad · Activity
Center · catalogues · value
reporting · training insights

Now Support

Standard case handling
stays here — Impact adds
enhanced P1/P2 targets,
Developer Support and
proactive engagements

AI stack

Now Assist / ServiceNow
Otto: outcome and
consumption summaries, AI
code fixes, GenAI root-cause
analysis

Business outcomes

Adoption · Platform health ·
Performance & availability ·
Value realised · Upgrade
readiness

2. Why ServiceNow introduced it

Five market drivers behind the product

- The platform outgrew ticket-based support. Customers stopped asking only “something is broken, fix it” and started asking “are we using what we bought?”
- Value realisation became a board-level question — finance wants ROI evidence, not anecdotes.
- Technical debt accumulated silently — leading practices were never measured continuously.
- Expert capacity is scarce — advice had to be packaged as fixed-scope delivery, not documents.
- Upgrade fatigue is real — the causes of broken upgrades needed to be attacked before the upgrade window.
- The product response is a closed loop: instrument → interpret → recommend → measure → repeat.

3. Architecture — components at a glance


What sits where, and what moves data

Layer	Component	Where it lives	Moves data off-instance?
Customer instance	Impact Store Application	Your instance — All → Impact	No
Customer instance	Impact Common / Content / Health	Dependent apps	No — findings stay local
Customer instance	Value Management data packs	Conditional per product; jobs must be enabled	Populates value metrics
Off-instance	Instance Observer	Off-instance cloud app (entitlement-tiered)	Telemetry collected off-instance
ServiceNow side	Impact Delivery Instance (IDI)	ServiceNow-hosted (impact.service-now.com)	Provider side
Integration	Service Bridge	Platform capability + dependent apps	Yes — bi-directional sync
Integration	Cloud Storage · Licensing Engine · MIF	Dependent apps	Yes — for the security review
AI	Now Assist / ServiceNow Otto skills	Your instance + Now Assist Admin	Governed by Now Assist terms

- Security review shortlist: Service Bridge, Cloud Storage, Licensing Engine and MIF move data. Impact Common, Content and Health do not.

3. Architecture — how Impact works

Instrument → interpret → recommend → measure → repeat



Foundations → Steady State cadence: monthly operational & health reviews · quarterly outcome, support and executive reviews

Callouts: Scan Engine findings are NOT transmitted to IDI · Feature availability depends on package, add-ons, roles and plugins

4. Feature catalogue — the five pillars

Each pillar maps to a client problem and a measurable output

Pillar	What it includes	The output the client sees
Platform Health	HealthScan definitions, Scan Engine, Real-Time Prevention, Proactive Code Check, AI fixes	Health scores by category, findings, remediation plans, upgrade readiness
Instance Observer	Performance and availability monitoring, alerts, triage, root-cause correlation, user experience	Alerts before users notice, faster MTTR, capacity evidence
Value Management	Objectives, outcomes, success metrics, Outcome Insights, Value Reports, Technical KPIs	Monetised value reporting with visible assumptions
Product Adoption	Capability Maps, Product Adoption Roadmaps, recommendations, Subscriptions	A sequenced adoption plan grounded in what we own
Experts & learning	Impact Squad, Developer Support, Accelerators, Initiatives, training insights, learning credits	Delivered artefacts and funded skills uplift

Platform Health — insight, prevention, remediation

Thousands of leading-practice checks, run continuously

- HealthScan definitions: the checks. Scan Engine: the runner. Findings: the output — and they stay in our instance.
- Persona dashboards mean each role sees what it owns: executive, platform owner, team lead, development team.
- Custom definitions let us encode our own standards; definition suites group them; exception reasons give an approved escape hatch.
- Real-Time Prevention validates changes as they are made rather than after the fact.
- Integration: findings can be routed into Jira, Azure DevOps, SPM or CWM so remediation lands in the team's existing backlog.
- AI: batch remediation with AI proposes fixes for review (needs Now Assist for Impact 3.03+, Now Assist for Platform 11.01+, the skill activated and `sn_impact_gen_ai.ai_fix.enabled = true`).
- First full scan can take hours — ServiceNow recommends running it overnight, especially in production.

Instance Observer — performance, availability, user experience

Monitoring from outside the instance — including when the instance is unwell

- Off-instance observability: if the instance degrades, its own monitoring is often the first casualty — Observer keeps working.
- Coverage includes core performance, database, ECC queue, email, events, host health, jobs, load balancer, node health, replicas, schedulers and semaphores.
- Alerting: popular and custom alerts, anomaly-based thresholds, RCC alerts and long-pending-job alert cards, with webhook notifications into our existing channels.
- Triage and root-cause correlation turn an alert into an investigation with evidence attached.
- User-experience insights require an admin to accept Terms & Conditions via Manage Permissions; user IDs are anonymised by default with an opt-in for non-anonymised collection.

Instance Observer — performance, availability, user experience (continued)

- Entitlement-tiered: seats, look-back window, alert counts and analytics differ between Guided and Total — verify. Not available in the mobile app.

Value management — proving ROI, not asserting it

Every outcome carries a metric, a baseline and a target

Caveat to state upfront

A meaningful Value Report needs at least two years of product data — start collecting early.

Element	What it is	Why it matters to the client
Business objective	High-level business statement, e.g. reduce cost to serve	Speaks the language of the executive sponsor, not IT
Business outcome	Measurable result connected to an objective	Makes the objective falsifiable
Success metric	The specific measure that moves	Removes ambiguity about what 'success' means
Reference quarter + goal	Baseline period and target	Enables trend reporting and comparison
Outcome Insights	On-track analysis against benchmark and goal	Tells us where to intervene, and recommends the capability or course
Value Reports	Monetised aggregation with visible assumptions	The artefact finance will challenge — so the assumptions are shown
Technical KPIs	Standardised technical measures with peer comparison	Answers 'how good are we?' with context

Product adoption and recommendations

Closing the gap between what we own and what we use

- Capability Maps show entitlements and adoption status, generated per instance — so we can see the gap between purchased and used.
- Product Adoption Roadmaps sequence capabilities into phases aligned to objectives, published and date-targeted.
- The recommendation engine surfaces the next best action on the Impact home page and in the IDI, each with rationale.
- Recommendations can be converted into work items in SPM or CWM (or Jira/Azure DevOps via integration) so they enter governed delivery.
- Accelerators and Initiatives are the delivery mechanism: fixed-scope expert engagements against the catalog.
- Concurrency governs how many can run at once — documented as 1 at a time for Guided and 7 concurrent for Total, plus Add-on concurrency.
- Subscriptions and the Consumption Report give transparent visibility of what was purchased and what was actually consumed.

Experts, learning and governance

The human layer that makes the tooling stick

Element	What it provides	How we consume it
Impact Squad	Customer Success Manager, plus CSE / Platform Architect / SAM in higher packages	Activity Center conversations, tasks, calendar, files
Developer Support	Named engineers who help troubleshoot our existing customisations	Named individuals raise cases; documented as 5 seats Guided / 10 Total — verify
Accelerators	Fixed-scope, expert-led engagements against four sub-catalogs	Requested from the app or IDI, consumed against concurrency
Initiatives	Multi-Accelerator programs towards a larger goal	Planned with the squad across quarters
Learning	Training Insights, contextual course recommendations, learning credits	Credits allocated against under-performing outcomes
Operating model	Foundations phase, then monthly and quarterly Steady State reviews	Calendarised governance with named attendees

5. Client scenarios — problem → capability → outcome (1/3)

Twelve worked scenarios from Australian enterprise and government programs

Problem	Impact capability	Measurable outcome
Renewal conversation you cannot support	Value Management: objectives → outcomes → success metrics → Value Reports	Defensible ROI at renewal — but start early; Value Reports need ~2 years of data
Shelf-ware we keep paying for	Capability Maps, Product Adoption Roadmaps, Subscriptions	Reduced shelf-ware, sequenced adoption instead of a wish list
Technical debt found too late	HealthScan, Scan Engine, update-set scans, Proactive Code Check, AI fixes	Fewer upgrade defects and a documented pre-upgrade health baseline
Performance reported by users, not monitoring	Instance Observer: anomaly alerts, triage, root-cause correlation	Detection before user reports; shorter MTTR with evidence attached

5. Client scenarios — problem → capability → outcome (2/3)

Problem	Impact capability	Measurable outcome
New platform owner inherits an undocumented program	Impact Squad, home page, Digest, Activity Center, roadmaps, training	Continuity of the value narrative within the first quarter
Change management is reactive	Real-Time Prevention, custom definitions and suites, exception reasons	Fewer change-related production incidents, with evidence of enforcement
Skills gap blocks adoption	Training Insights, contextual learning, learning credits	Measurable improvement on the linked outcome metric
Remediation work never gets prioritised	Recommendation → work item conversion (SPM/CWM/Jira/ADO), Accelerators	Rising remediation completion; no dual-tracking of platform work

5. Client scenarios — problem → capability → outcome (3/3)

Problem	Impact capability	Measurable outcome
Multiple business units, competing priorities	Stakeholder Groups (Group Views), per-group objectives, instance filters	Higher executive engagement; cleaner value attribution per function
Controlled environment, constrained feature set	Documented availability restrictions; squad guidance; no domain separation	No late-stage design rework from an entitlement or environment constraint
Consumption and entitlement governance	Consumption Report / Benefits & Usage, Accelerator and Developer Support usage	Higher entitlement utilisation; evidence-based position at renewal
Regulated onboarding: manual registration	Manual registration path, Service Bridge, connection verification	Compliant verified connection showing Active replication — the gate for everything else

6. Business benefits by stakeholder

What each role gets — and the evidence they can point to

Stakeholder	What they get	Evidence
CIO / CDIO	Confidence the platform delivers intended outcomes; defensible ROI position	Value Reports, Outcome Insights, quarterly executive reviews
CFO / Finance	Visible return on existing spend; reduced shelf-ware; assumption-transparent monetisation	Value Reports, Capability Maps, Consumption Report
CIO office / Strategy	Prioritised roadmap aligned to objectives; recommendations converted to governed work	Roadmaps, SPM/CWM work items, Accelerator outcomes
Platform Owner	Continuous health and performance visibility; defined cadence; expert escalation path	Health scores, Observer alerts, Monthly Health Assessment, Tech KPIs
Platform Architect	Engineering-quality evidence for design decisions; upgrade readiness	Scan findings, custom definitions, Proactive Code Check, peer KPI comparison
Service Ops Manager	Fewer user-reported incidents; faster resolution of hard cases	Alert-to-resolution history, Developer Support outcomes
Developers	Specific actionable findings; fast paths to fixes	Team Lead dashboards, AI fix proposals, Jira/ADO integration
Security & Risk	Documented data flows, consent controls, environment limitations	Consent records, data-egress component list, availability restrictions

Business benefits by category

Eight lenses for the business case

Financial

Return on existing
licence spend; avoided
re-purchase;
monetised value
reporting

Risk reduction

Technical debt
measured
continuously; upgrade
and change risk found
pre-production

Operational efficiency

Remediation routed
into existing backlogs;
faster triage and root
cause

Adoption & enablement

Capability maps and
roadmaps turn
entitlements into used
capability

Governance & assurance

Foundations → Steady
State cadence with
auditable artefacts

Decision quality

Recommendations
with rationale, peer
context and expert
validation

Predictability

Managed upgrade
readiness instead of
emergency response

Scalability & continuity

Expert capacity
without permanent
headcount; knowledge
retained in the
platform

7. Impact versus traditional support

Reactive assurance versus proactive value creation — and how they coexist

Dimension	Traditional reactive support	ServiceNow Impact
Trigger	A case is raised after something breaks	Continuous instrumentation — scans, telemetry, value data
Scope	Break-fix of ServiceNow base functionality	Plus health, performance, adoption, value, upgrade readiness
Deliverable	Fix, workaround, known error	Assessments, roadmaps, monetised value reports, expert-led Accelerators
Expert access	Support agents via Now Support portal	Named Impact Squad, plus Developer Support engineers
Cadence	Ad hoc, incident-driven	Monthly operational and health reviews; quarterly outcome and executive reviews
Success measure	SLA adherence and time to close	Outcomes achieved, value realised, health trend, adoption
Commercial model	Included with product support entitlements	Subscription packages + Add-on SKUs + Accelerator concurrency

- What Impact does not do: it does not remove the need for administrators, change control, data governance, or a partner for net-new implementations.

8. Prerequisites — readiness before we start

Most stalled deployments stall here, not on technology

- Platform: Yokohama or Xanadu Patch 2 and above; Impact Store Application entitlement (sn_entitlement) version 4.1.2 or higher — verify the live Store listing.
- Commercial: confirm the package (Guided / Total / Integrated Success) and any Add-ons; entitlements differ materially.
- People: a named contact administrator to receive the registration email, a Platform Owner, an Impact App Admin and a Scan Engine Admin.
- Access: Impact Squad access approval (read-only, 30 days, revocable, quarterly re-approval).
- Consent: Terms & Conditions acceptance for user-experience monitoring; consent when scheduling the Monthly Health Assessment.
- Security: dedicated integration user, authentication method chosen (OAuth 2.0 preferred over Basic), data-egress review of Service Bridge, Cloud Storage, Licensing Engine and MIF.
- Data: exact instance names matching the instance_name property, full URLs, one Production designation.
- Capacity: storage growth (attachments dominate) and a cutover window for the first full scan.
- Known limitation: Impact does not support domain separation; portions are unavailable in restricted environments including Australia IRAP-Protected.

8. Installation and activation — the sequence (1 of 2)

Each Guided Setup section must be completed to unlock the next

#	Step	Navigation	Role
1	Procure the app	ServiceNow Store → Apps and Solutions → search Impact → Buy	Any Impact role
2	Install the app	All → Application Manager → search Impact → Install (Sync now if missing)	admin
3	Open Guided Setup	Application Manager → Configure, or All → Impact → Guided Setup → Get Started	Impact App Admin / admin
4	Onboard users	All → Impact → Configuration → Guided Setup → User management	Impact App Admin
5	Platform Health users & teams	All → Impact → Guided Setup → Assign Platform Health users; then Scan Engine Properties → Team Leads → New	Impact App Admin / admin
6	Activate Scan Engine	Guided Setup → Impact Platform Health → Activate Scan Engine; grant Can read on sys_update_version	Impact App Admin / admin
7	(Optional) AI code fixes	Scan Engine Properties → Real Time Scanning → Enforce real-time validation = true; Now Assist → Skills → Impact → Code Fix → Activate skill; sn_impact_gen_ai.ai_fix.enabled = true	Scan Engine Admin
8	Run the first full scan	All → Impact → Platform Health → Scheduled Scan → Execute Now; monitor Scan Engine → Scan Status	Impact App Admin / admin

- Expect the first full scan to take hours — run it overnight, especially in production. Scan Engine findings are not transmitted to the Impact Delivery Instance.

8. Installation and activation — the sequence (2 of 2)

Connection, registration, migration, then the value layer

#	Step	Navigation	Role
9	Register instances	ALL → Impact → Configuration → Scan Engine Properties → My SN Instances → New	sn_se.scan_engine_admin
10	Configure authentication	OAuth 2.0 preferred (production or development variant) or Basic authentication	admin / Scan Engine Admin
11	Validate connection	My SN Instances → open record → Validate Connection → status 'Connection valid'	sn_se.scan_engine_admin
12	Automated registration	All → Impact → Configuration → Guided Setup → Register your instance	Impact App Admin + IDI admin
12b	Manual registration (regulated / GCC)	IDI → Activity Center → Instance registration → Add instance	admin / any Impact role
13	Verify data connection	Guided Setup → Verify the Connection — success is Inbound + Outbound 'Active replication'	Impact App Admin
14	Initiate sync & migration	All → Impact → Configuration → Guided Setup → Initiate sync & migration → Start Data Migration	Impact App Admin / admin
15	Enable value layer	Objectives and outcomes; per-product data collection jobs (dc-*-install / activate / config); Instance Observer configuration	Impact Admin / Platform Owner

- Validation warning: resolve the root cause before retrying a failed connection — repeated failures trigger the account lockout policy. Two flags must be cleared to recover: locked_out and password_needs_reset.

9. The first 90 days — Foundations

From activation to the first credible read-out

- Week 0 — kickoff and alignment: Impact Kickoff Meeting, success criteria agreed, named contacts confirmed, squad access approved.
- Weeks 1-2 — technical enablement: app installed, Guided Setup started, users and roles assigned, Platform Health users placed in groups, development teams created.
- Weeks 2-3 — baseline the platform: Scan Engine activated, first full scan run overnight, findings triaged, remediation backlog created.
- Weeks 2-4 — connect and migrate: instances registered, authentication configured, connection validated, registration completed, Active replication confirmed, data migration run.
- Weeks 3-6 — agree the value story: objectives and outcomes agreed with baselines and targets; data collection jobs enabled per product.
- Weeks 4-8 — plan the work: Capability Map reviewed, Product Adoption Roadmap published, recommendations converted into governed work items.
- Weeks 6-10 — extend observability: Instance Observer enabled, thresholds tuned, webhook notifications routed, consent captured if in scope.
- Weeks 8-12 — first value read-out: Monthly Health Assessment reviewed, outcome trends analysed, learning credits allocated.

9. The first 90 days — Foundations (continued)

- Exit criteria: connection verified, full scan complete with owned findings, outcomes with baselines, published roadmap, calendarised reviews.

Steady state — the repeating cadence

What happens monthly, quarterly and annually once Foundations is complete

Cadence	Forum	Inputs	Outputs
Monthly	Operational and health review (Platform Owner + CSM + team leads)	New scan findings, health score movement, Observer alerts, remediation backlog	Prioritised actions with owners and dates
Monthly	Health assessment review	Monthly Health Assessment output (reports typically ~2 weeks after the assessment)	Accepted recommendations, next month's focus
Quarterly	Outcome review (business stakeholders + CSM)	Outcome Insights vs benchmark and goal, adoption status, roadmap progress	Roadmap adjustments, new or retired outcomes, Accelerator re-planning
Quarterly	Support and executive review	Case trends P1-P4, Developer Support usage, Accelerator consumption, learning balance	Value narrative update, investment decisions
Quarterly	Governance hygiene	Squad access (read-only, 30 days, revocable)	Renewed or revoked access
Annually	Value realisation	Value Reports with assumptions, Consumption Report	Renewal position, package right-sizing, next-year objectives

Verify before you commit

Entitlement-dependent or documentation-inconsistent points to confirm first

Point to verify	Why
Package deltas (Guided / Total / Integrated Success)	Custom definition limits, Observer seats and look-back, alert counts, Accelerator concurrency and Developer Support seats all differ
P1 / P2 response targets by package	Documented as 30/120 minutes (Guided) and 15/60 minutes (Total)
Proactive Code Check status	Documented as deprecating from Impact Zurich 6.0.8 — confirm before multi-year planning
Group Views availability	Documented as initially available to limited customers; requires squad enablement
On-demand Accelerators	Documented as not available to Guided — confirm for our package
Restricted-environment availability	Portions of Impact are unavailable in FedRAMP, NSC DOD IL5 and Australia IRAP-Protected environments, and for self-hosted users and MSPs (except internal use)
AI capabilities	Now Assist / ServiceNow Otto skills require version, entitlement and activation (e.g. Now Assist for Impact 3.03+, Now Assist for Platform 11.01+)
Domain separation	Not supported for Impact — a design constraint, not a configuration option
Value Reports	Need at least two years of product data to be meaningful

Next steps

What we are asking for

1 • Confirm entitlement

Verify the exact package, Add-ons and entitlement figures in writing with ServiceNow before committing to a plan.

2 • Name the people

Named contact administrator, Platform Owner, Impact App Admin and Scan Engine Admin.

3 • Approve access & consent

Squad access (read-only, 30 days, quarterly re-approval) and the consent decisions for monitoring.

4 • Schedule the sequence

Procure → install → Guided Setup → scan (overnight) → register → validate → Active replication → migrate.

5 • Agree the value story

Objectives and outcomes with metrics, baselines and targets — before the first report is expected.

6 • Calendarise governance

Lock the monthly operational/health review and the quarterly outcome and executive reviews.

Summary — the client message

Five sentences to remember

- Impact is a subscription customer success product: a workspace in our instance, an expert workspace at ServiceNow, and named people — not another module.
- It answers four questions: is the instance healthy, is it performing, are we realising value, and what should we do next.
- Its evidence is artefacts we can show: health scores, findings, adoption roadmaps, and monetised value reports with visible assumptions.
- It coexists with Now Support rather than replacing it, and it does not remove the need for our own administration and governance.
- About half the value depends on foundational work we control — connection, roles, scans and data collection — so the sequence matters more than the licence.

Appendix · Glossary

Key Impact terms used in this briefing

Term	Meaning
Impact Store Application	The Impact application installed in your ServiceNow instance — the consumer-side UI (All → Impact)
Impact Delivery Instance (IDI)	The ServiceNow-hosted provider instance where the Impact Squad works (impact.service-now.com)
Service Bridge	The integration technology enabling secure bi-directional sync between your instance and the IDI
HealthScan	The library of leading-practice checks the Scan Engine runs against your instance
Scan Engine	The component that executes HealthScan definitions and produces findings in your instance
Finding	A specific scan result — the unit of remediation work; findings are not sent to the IDI
Capability Map	A per-instance view of entitled capabilities and their adoption status
Product Adoption Roadmap (PAR)	A phased, published plan sequencing capabilities towards objectives
Outcome	A measurable business result linked to an objective, with a success metric, baseline and goal
Value Report	A monetised aggregation of outcome performance with visible assumptions
Accelerator / Initiative	A fixed-scope expert-led engagement / a programme sequencing several Accelerators
Concurrency	The number of Accelerators that may run at the same time under your package
Instance Observer	Off-instance performance and availability monitoring (also written 'Instance Observatory')
ServiceNow Otto	The Australia-release rebrand of Now Assist capabilities in this context
CIP	Customer Impact Plan — Impact's planning construct with the squad

Appendix · Decision points and FAQs

The questions that usually come up next

Question	Short answer
Does Impact replace our support contract?	No. Now Support continues to handle cases; Impact adds enhanced P1/P2 targets, Developer Support and proactive engagements.
Will it fix our technical debt?	It finds, prioritises and routes it, and can propose AI-assisted fixes. Remediation still consumes our delivery capacity.
When will we see hard ROI?	Value Reports need at least two years of product data, so outcomes must be baselined early to report later.
Do we need to be on a particular release?	Yokohama or Xanadu Patch 2 and above, with sn_entitlement 4.1.2+ — verify the live Store listing.
Does our data leave the instance?	Service Bridge, Cloud Storage, Licensing Engine and MIF move data. Impact Common, Content and Health do not; Scan Engine findings stay local.
Can we use it with domain separation?	No — domain separation is not supported for Impact.
What about IRAP-Protected environments?	Portions of Impact are documented as unavailable in restricted environments including Australia IRAP-Protected data centres — confirm the feature subset.
How many Accelerators can run at once?	Documented as 1 at a time for Guided and 7 concurrent for Total, plus Add-on concurrency — verify.
Who runs the Monthly Health Assessment?	ServiceNow delivers the assessment; ServiceNow is not responsible for implementing or managing it — that ownership is ours.
What is the first thing to do?	Name the contact administrator, confirm entitlement, then procure and install the Store Application.

Speaker notes — Slide 1: ServiceNow Impact

Opening frame.

Purpose of today: agree what Impact is, what it is not, what it changes for us operationally, and what we would need to do to activate it.

Three things to land by the end: (1) Impact is a subscription product that adds a measurement and expert layer on top of the platform we already run — it does not replace administration or Now Support; (2) roughly half of its value depends on foundational work we control — instance connection, roles, scans and data collection; (3) several capabilities are entitlement-dependent, so nothing in this deck should be treated as committed until we confirm our package.

Ask for the session: agreement on the activation sequence and an owner for each prerequisite.

Speaker notes — Slide 2: Agenda

Set expectations on depth. Sections 1–3 are the conceptual frame, 4–6 are the value case, 7 is the honest comparison, 8 is the operational ask.

Everything in the deck is drawn from the Australia-release product documentation listed in the written brief (Appendix E). Where the documentation is silent or a figure is entitlement-dependent, the slide says “verify” rather than guessing.

If the audience is purely executive, the recommended path is slides 1–7, then 17–19, then 23–26. The technical sections (9–16, 20–22) are for the architecture and platform audience.

Speaker notes — Slide 3: Executive summary

The three-line version.

1. Prove value. Impact's value layer connects business objectives to measurable outcomes with a metric, a baseline quarter and a target — and rolls that into monetised Value Reports where the assumptions are visible. Caveat to state upfront: a meaningful Value Report needs at least two years of product data, so the earlier we start collecting, the earlier we can report.
2. Know your health. Thousands of leading-practice checks run against the instance continuously rather than being discovered during an upgrade. Findings stay in our instance. Instance Observer watches performance and availability from outside the instance.
3. Plan and execute. Capability maps show what we own versus what we actually use; roadmaps sequence the next phases; Accelerators are fixed-scope expert engagements consumed against package concurrency.

The honest framing: Impact gives us measurement, guidance and expert capacity. It does not do our remediation for us.

Speaker notes — Slide 4: 1. What ServiceNow Impact is

Define it precisely, because “Impact” is often heard as “support” or “another module”.

It is a subscription product with three delivery surfaces. The one that matters most for day-to-day work is the Impact Store Application inside our instance — reached at All → Impact. The IDI is where our named squad works; we reach it through the Activity Center. Human expertise is the third piece: CSM, and in higher packages a Customer Success Executive, Platform Architect and Support Account Manager.

The directional statement worth emphasising: from Yokohama onward, new innovative Impact features land in the Store Application, not the IDI. So the app is where our future capability will accumulate.

Close with the boundary: this is not a replacement for our admins or our support contract, and we should not describe it that way to the business.

Speaker notes — Slide 5: Why it exists and what it solves

This slide reframes the conversation from support to outcomes.

Walk the rows rather than reading them. The pattern is that each row is a question our own stakeholders already ask — usually at renewal, or after an incident, or when a business unit questions the spend.

Two rows are worth dwelling on with an executive audience. First, “cannot prove value” — that is the value management row and it is the reason many organisations adopt Impact at all. Second, “what should we do next” — most stalled programs are not short of findings; they are short of a defensible priority order.

For a technical audience, the two rows that matter are the health and performance rows — those are the ones that change day-to-day engineering work.

Speaker notes — Slide 6: Where Impact fits on the platform

Two rows, three boxes each — it reads as “what we run”, “what connects us”, “what ServiceNow runs for us”, then where support and AI sit.

Key point for the platform team: the Service Bridge connection must reach Active replication in both directions. That single status is the gate for almost everything else in the product, and it is the most common place where activation stalls.

Key point for the support/operations audience: Now Support does not go away. Case handling stays there; Impact adds response targets and Developer Support on top.

Key point for everyone: impact sits around and above the platform. It does not replace our change, release or data governance processes.

Speaker notes — Slide 7: 2. Why ServiceNow introduced it

This is the “why does this product exist” slide. It is useful when a client is sceptical that Impact is more than rebranded support.

The strongest of the five with an Australian enterprise or government audience is usually value realisation — audit and budget scrutiny make the ROI question unavoidable.

For a technical audience, lead with upgrade fatigue and silent technical debt: those are the pain points the platform team already feels weekly.

Land the closed loop at the end because it is the mental model for the architecture slide that follows.

Speaker notes — Slide 8: 3. Architecture — components at a glance

This is the slide the security and architecture reviewers will photograph.

The critical answer to the inevitable question “what leaves our instance?” is in the last column. Four components move data: the Service Bridge stack, Cloud Storage, the Licensing Engine and the MIF Customer Instance. Impact Common, Impact Content and Impact Health do not move data, and Scan Engine findings stay in our instance.

Note the Instance Observer row: it is off-instance monitoring by design — that is precisely why it still works when the instance itself is degraded, and it is also why it needs its own consent conversation for user-experience data.

Values, licensing and subscription data feed Capability Maps and the Subscriptions view — worth naming so nobody is surprised later.

Speaker notes — Slide 9: 3. Architecture — how Impact works

The one diagram to walk slowly.

Left to right: our estate, the secure sync layer, the ServiceNow side, then the business outcomes. Underneath, the operating cadence that closes the loop.

Column 1 is everything under our control and our responsibility: the app, the Scan Engine and HealthScan definitions, the Performance Analytics collection jobs that feed value metrics, our users and roles, and off-instance observability.

Column 2 is the single integration point. Automated registration is preferred; regulated and GCC customers must use manual registration.

Column 3 is what ServiceNow brings: named squad, the Activity Center, value management, and the catalogues.

Column 4 is why we are doing this.

Two callouts to say out loud: Scan Engine findings are NOT transmitted to the IDI; and feature availability depends on our package, add-ons, roles and installed plugins — which is why we verify before we commit.

Speaker notes — Slide 10: 4. Feature catalogue — the five pillars

This is the map for the rest of the feature discussion.

The framing to use: every pillar exists to answer a question a stakeholder already has. Platform Health answers “is the instance in good shape?”. Instance Observer answers “is it performing?/will users notice?”. Value Management answers “was it worth it?”. Product Adoption answers “what should we do next?”. Experts and learning answer “who helps us, and how do we get better?”

If time is short, the two pillars to cover properly are Platform Health (because it is the most immediately useful to the platform team) and Value Management (because it is the one the CFO will ask about).

The next four slides go one level deeper; the detailed per-feature structure is in the written brief, section 3.

Speaker notes — Slide 11: Platform Health — insight, prevention, remediation

This is the pillar the platform and development teams will care about most.

Start with the locality point because it is a common objection: findings are produced in our instance and are not transmitted to the Impact Delivery Instance.

The persona dashboards matter more than they sound — they are what stop findings becoming an undifferentiated list. A team lead sees their team's new findings after a sprint; an executive sees the trend.

On custom definitions: this is where we encode house rules (no business logic in UI policies, no direct CMDB writes from integrations). Guided has a documented cap of up to 10 active custom definitions and Total is documented as unlimited — verify against our package.

On the first scan: set expectations now. It can take hours; run it overnight in production and monitor Scan Status rather than re-triggering.

Speaker notes — Slide 12: Instance Observer — performance, availability, use

Position this as the answer to the most demoralising class of incident: performance problems reported by users after the fact, with no evidence left to investigate.

The off-instance argument is the headline: when an instance is degraded, in-instance monitoring is unreliable. Observer watches from outside.

For the operations audience, the practical wins are anomaly-based alerting (fewer noisy static thresholds) and webhook routing into the existing on-call tooling rather than a new console nobody watches.

Raise the consent point deliberately: user-experience monitoring requires Terms & Conditions acceptance by an admin in Manage Permissions, IDs are anonymised by default, and there is an opt-in for non-anonymised collection. That is a conversation for our privacy and employee-representation stakeholders, not a technical toggle.

Finally, flag the entitlement differences and the mobile-app gap so expectations are set correctly.

Speaker notes — Slide 13: Value management — proving ROI, not asserting it

This is the pillar that justifies the purchase to the CFO, and the one most often run badly.

The discipline is in the chain: objective → outcome → success metric → reference quarter and goal. If any link is vague, the report is not defensible. Push the sponsor to make objectives business statements (“reduce cost to serve”) rather than IT statements (“complete the next upgrade”).

Outcome Insights is the operational half — it tells us which outcomes are off track and recommends the capability or the training that addresses them. Value Reports are the reporting half, and their credibility comes from showing assumptions rather than hiding them.

State the caveat out loud rather than letting someone discover it later: a meaningful Value Report needs at least two years of product data. That is an argument for starting data collection now, even if the first report is modest.

Speaker notes — Slide 14: Product adoption and recommendations

Three connected ideas: see the gap, plan the closure, execute with help.

The gap point is usually the most confronting for clients — most organisations materially under-use what they already own, and that is a renewal risk as much as an efficiency one.

The planning point is that roadmaps are sequenced against business objectives, not product release order. That is what makes them adoptable by the business.

The execution point is the conversion step: a recommendation that never becomes a work item is just a notification. Route them into SPM, CWM, Jira or Azure DevOps during normal planning.

Raise concurrency honestly: it is a real constraint on how many expert engagements can run at once, and scheduling is subject to availability. Plan engagements in advance rather than assuming immediate start.

Speaker notes — Slide 15: Experts, learning and governance

Tooling without a human layer produces dashboards nobody acts on — this slide is the counterweight to the product slides.

The squad is the primary relationship: the CSM orchestrates onboarding, the Foundations phase and the review cadence.

Be precise about Developer Support scope, because it is commonly over-estimated: it covers troubleshooting and improving existing customisations, and explicitly does not cover net-new implementations or customisations that break after a family upgrade — those go to Now Support.

Learning is the most commonly wasted entitlement because credits expire unused. The mechanism that makes it purposeful is the link from an under-performing outcome to a specific recommended course.

Close on governance: Impact works when it is calendarised. That is the ask in section 8.

Speaker notes — Slide 16: 5. Client scenarios — problem → capability → outcome

Scenario group 1 of 3. Use these as discussion prompts rather than a script.

Ask which of these the client recognises as their own situation, then go deep on that one. The strongest openers for most organisations are the renewal-value scenario and the technical-debt scenario.

Each scenario is deliberately written as problem → capability → measurable outcome so the conversation cannot drift into feature description without landing on a result.

Speaker notes — Slide 17: 5. Client scenarios — problem → capability → outcome

Scenario group 2 of 3. Use these as discussion prompts rather than a script.

Ask which of these the client recognises as their own situation, then go deep on that one. The strongest openers for most organisations are the renewal-value scenario and the technical-debt scenario.

Each scenario is deliberately written as problem → capability → measurable outcome so the conversation cannot drift into feature description without landing on a result.

Speaker notes — Slide 18: 5. Client scenarios — problem → capability → outcome

Scenario group 3 of 3. Use these as discussion prompts rather than a script.

Ask which of these the client recognises as their own situation, then go deep on that one. The strongest openers for most organisations are the renewal-value scenario and the technical-debt scenario.

Each scenario is deliberately written as problem → capability → measurable outcome so the conversation cannot drift into feature description without landing on a result.

Speaker notes — Slide 19: 6. Business benefits by stakeholder

Use this slide to make the benefits land with the person who owns the budget or the risk.

The column that does the work is the last one — evidence. Benefits without evidence get discounted in the room; naming the artefact that proves it is what makes the benefit credible.

If the audience is mixed, pick two rows: CIO for the value story and Security for the constraints story. Those are usually the two competing narratives that decide whether the program proceeds.

The Security row is deliberate: naming the constraints and the controls openly is what gets an approval rather than a deferral.

Speaker notes — Slide 20: Business benefits by category

Eight short lenses; do not read all eight aloud. Pick the three that match the client's stated priorities.

For cost-pressured organisations: financial, operational efficiency and predictability.

For regulated or high-assurance organisations: risk reduction, governance and assurance, and decision quality.

For organisations in the middle of a transformation: adoption and enablement, and scalability and continuity — because the constraint there is almost always skilled people, not licences.

The honest note worth adding verbally: benefits 1–3 depend on us acting on findings; benefits 4–8 depend on us keeping the cadence. Neither is automatic.

Speaker notes — Slide 21: 7. Impact versus traditional support

Handle this slide honestly — it is where scepticism is highest and where over-claiming does the most damage.

The correct framing is coexistence, not replacement. Now Support remains the home of incident management; Impact adds enhanced P1/P2 response targets (documented as 30/120 minutes for Guided and 15/60 for Total — verify), Developer Support, and the proactive layer.

Read the “what Impact does not do” line aloud. It buys credibility for everything else on the slide, and it prevents the two most common client misapprehensions: that Impact replaces support, and that Impact will fix technical debt by itself.

The most useful single sentence: support can tell us the platform is healthy today; Impact tells us whether the platform is delivering what the business paid for.

Speaker notes — Slide 22: 8. Prerequisites — readiness before we start

This slide is the operational ask. Do not rush it.

The single most important line for scheduling is the named contact administrator — the registration email goes to that person, and if they are not identified (and available), the whole sequence waits.

The second is the Security item: OAuth 2.0 is preferred over Basic authentication, and the data-egress review should be scoped to the four moving components rather than the whole product. Doing that early prevents a late security stall.

State the two limitations plainly: no domain separation, and restricted-environment unavailability including Australia IRAP-Protected. If either applies to us, we confirm the exact feature subset with the squad before design sign-off.

The pre-implementation checklist with owners is in section 7.10 of the written brief — that is the artefact to take away from this slide.

Speaker notes — Slide 23: 8. Installation and activation — the sequence (1 of

Steps 1-8 are everything up to the point where the instance is scanning.

Two sequencing traps worth calling out. First, Guided Setup steps lock — you cannot jump ahead, and marking a step complete prematurely hides a failure until much later in the sequence. Mark complete only on genuine success.

Second, step 6 has a non-obvious sub-step that blocks most first installations: when the Scan Engine error banner appears, you must follow the `sys_update_version` link to the Tables page, open the Application Access tab, select Can read and update. Without that, scanning does not start.

Step 8 timing: a first full scan can legitimately take hours on a large instance. Run it overnight in production and watch Scan Status rather than re-triggering it.

Speaker notes — Slide 24: 8. Installation and activation — the sequence (2 of

Steps 9–15 are the connection and value-enablement half.

Step 9 has three data-quality requirements that cause most validation failures: the instance name must match the instance_name system property exactly, the URL must be exactly `https://instancename.service-now.com` with no `www.`, no trailing slash and no path, and only one instance in the stack may be designated Production.

Step 11 is where impatience is expensive. Every failed validation attempt contributes to the account lockout policy. If validation fails, fix the root cause first — do not retry blind. If the account does lock, remember that two flags must be cleared, `locked_out` and `password_needs_reset`; clearing only one is insufficient.

Step 15 is the step people forget. Installing a data collection pack does not collect data — the jobs must be enabled. And objectives and outcomes should be agreed early, because Value Reports need at least two years of product data.

Speaker notes — Slide 25: 9. The first 90 days — Foundations

Foundations is roughly the first 90-120 days and its job is to get the loop running end to end and produce a first credible read-out.

Two things to emphasise. First, the ordering is deliberate: connection and scans before value reporting, because the value layer needs data to accumulate. Second, the milestone that matters most is at week 12 — not “everything installed” but “findings owned, outcomes baselined, roadmap published, reviews in diaries”.

The most commonly skipped item is the last exit criterion: calendarising the reviews. Programs that do not calendarise the cadence quietly stop using the product within two quarters.

Useful framing for the sponsor: by day 90 we should be able to show the health baseline and the first outcome trends, even though a full value report needs considerably more data history.

Speaker notes — Slide 26: Steady state — the repeating cadence

Steady State is a repeating cadence, not a project phase — this slide is the operating contract.

The two meetings that carry the most weight are the monthly operational and health review (which keeps technical work moving) and the quarterly executive review (which keeps sponsorship alive). If only one can be protected, protect the quarterly executive review.

Note the practical detail on health assessment reporting: reports typically arrive around two weeks after the assessment, so plan the review meeting accordingly rather than expecting same-week output.

Also note what is explicitly not ServiceNow's responsibility: implementation and management of the Monthly Health Assessment sits with us. Budget internal ownership for it rather than assuming the squad runs it end to end.

Speaker notes — Slide 27: Verify before you commit

This is the integrity slide, and it is worth its place in the deck.

The single most common source of client disappointment in Impact programs is assuming a capability is included when it is entitlement-dependent. Going through this list openly — and committing to confirm each item against the contract — is what prevents that.

Three items deserve specific mention. Domain separation is simply not supported, so if our design assumes it, we need a different design. Restricted-environment availability matters directly for Australian government and defence-adjacent estates. And Value Reports need at least two years of data — that is an argument for starting now, not a reason to defer.

The message to land: everything on this slide is a question we will answer in writing before we commit to a plan.

Speaker notes — Slide 28: Next steps

Close with a specific, ownable ask rather than a summary.

These six items map directly onto the prerequisites and activation sequence, and each one has a named owner in the pre-implementation checklist (section 7.10 of the written brief).

If the client can only commit to some of them, the two that matter most for schedule are items 2 and 4: naming the contact administrator, and booking the first scan's cutover window. Everything else can follow, but those two gate the start.

Offer the written brief and the pre-implementation checklist as the take-away, and propose the follow-up session as the entitlement confirmation rather than a general review.

Speaker notes — Slide 29: Summary — the client message

This is the closing slide; keep it to the five sentences.

If the audience remembers nothing else: Impact is a measurement and expert-capacity layer, its value is demonstrated through artefacts rather than assertions, and it requires foundational work from us before it can deliver anything.

Invite the two questions that usually decide the next step: “what exactly is in our package?” and “who owns the activation sequence internally?” Both are answered in the written brief.

Speaker notes — Slide 30: Appendix · Glossary

Reference slide — do not present it line by line. It exists so the deck is self-contained when it is circulated without a presenter.

If a question arises mid-session on terminology, this is the place to jump to. The three terms most often confused are Impact Store Application (in our instance) versus Impact Delivery Instance (ServiceNow-hosted), Accelerator versus Initiative, and concurrency.

One terminology warning worth stating if “Success Plans” comes up: that name belongs to ServiceNow's separate Customer Success Management product, not Impact. Inside Impact the planning constructs are the Customer Impact Plan, Product Adoption Roadmaps and objectives/outcomes.

Speaker notes — Slide 31: Appendix · Decision points and FAQs

Anticipated-question slide for the Q&A and for circulation.

The three questions that decide whether a program proceeds are usually the first three: does it replace support, will it fix our debt, and when do we see ROI. Each answer here is deliberately calibrated to under-promise rather than over-claim.

The data-egress question will come from security, and the domain-separation and IRAP questions from architecture or risk. Having them answered in the deck saves a follow-up session.

Anything not answered here is addressed in the written brief, which carries the full source-topic list in Appendix E.